

# Data Mining 2

## Topic 04 : Hyperparameter Tuning

### Lecture 02 : Hyperparameter Tuning

Dr Kieran Murphy

Department of Computing and Mathematics, Waterford Institute of Technology.  
(Kieran.Murphy@setu.ie)

Spring Semester, 2026

#### Outline

- Standard Techniques - Grid vs Random Search
- Bayesian optimisation - hyperopt and Scikit-Optimize

## Recap of where we are (were from previous lecture)

- Given a pipeline / model we “can” tune each of the hyperparameters by constructing a validation curve.
- Problems ... (manual approach is not practical anything but small problems)
  - Can have a huge (100s – 1,000s) number of parameters.
  - A sequence of 1-dimensional searches is not the same as one  $d$ -dimensional search — interplay between hyper-parameters.
- Approaches / Techniques ...
  - **Grid Search** — Systematic, regular, predetermined deterministic sample of parameter space.
  - **Random Search** — Non-adaptive, random sample of parameter space.
  - **Bayesian Search** — Adaptive, random sample of parameter space.

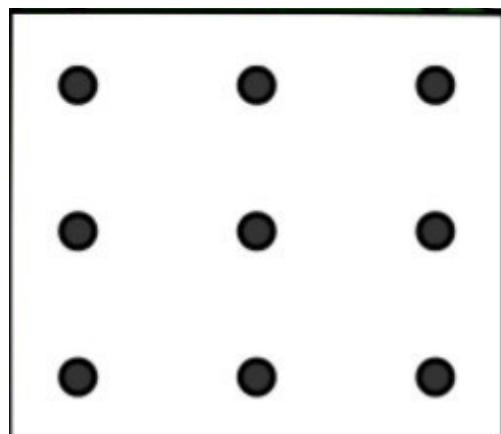
# Grid Search

Validation curves are important but examining at each parameter individually is time consuming and does not take into account interactions between parameters — in effect it is a series of 1D semi-manual searches.

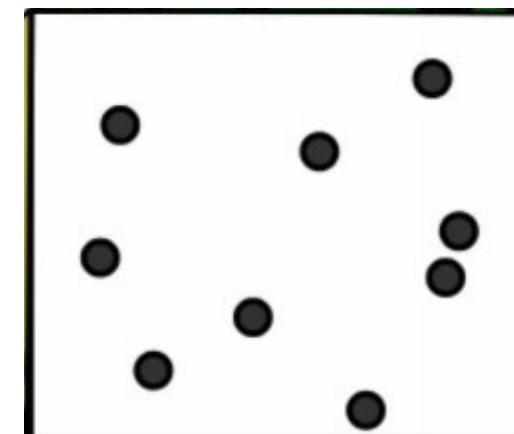
## Grid Search

Try all available parameters combinations, 1 by 1, and choose the one with the best cross validation results.

Grid Search



Random Search



## Grid Search (`sklearn.model_selection.GridSearchCV`)

### Grid Search

Try all available parameters combinations, 1 by 1, and choose the one with the best cross validation results.

#### Drawbacks

- Still very slow — trying ALL combinations of ALL parameters with no modification of search based on results found to date.
- ALL combinations  $\implies$  every additional hyperparameter to vary multiplies the number of iterations you need to complete.
- It can work only with discrete values.

If the global optimum is on `n_estimators=550`, but you are doing `GridSearchCV` from 100 to 1000 with step 100, you will never reach the optimal point.

#### Strategies

- You need know / guess the approximate region of the optimum to start.
- To mitigate against drawbacks, do multiple lower dimensional grid searches, or repeat search with narrower grids and smaller step sizes.

# Grid Search Example 1

```
1 from sklearn.svm import SVC
from sklearn.model_selection import GridSearchCV
pipeline = Pipeline([
    ('scl', StandardScaler()),
    ('clf', SVC(random_state=SEED))])

param_range = np.logspace(-4, 4, 10)
param_grid = [{}
    'clf__gamma': param_range,
    'clf__C': param_range,
    'clf__kernel': ['rbf']
}]
gs = GridSearchCV(estimator=pipeline,
    return_train_score=True,
    param_grid=param_grid, scoring='accuracy', cv=10, n_jobs=-1)
```

- Switched from LogisticRegression to SVM — more hyperparameters to tweak.
- How many parameter combinations were generated/used?

# Grid Search Example 1

To perform grid search, we call the `fit` method as usual ...

2  
2. gs = gs.fit(X\_train, y\_train)  
print(gs.best\_score\_)  
print(gs.best\_params\_)

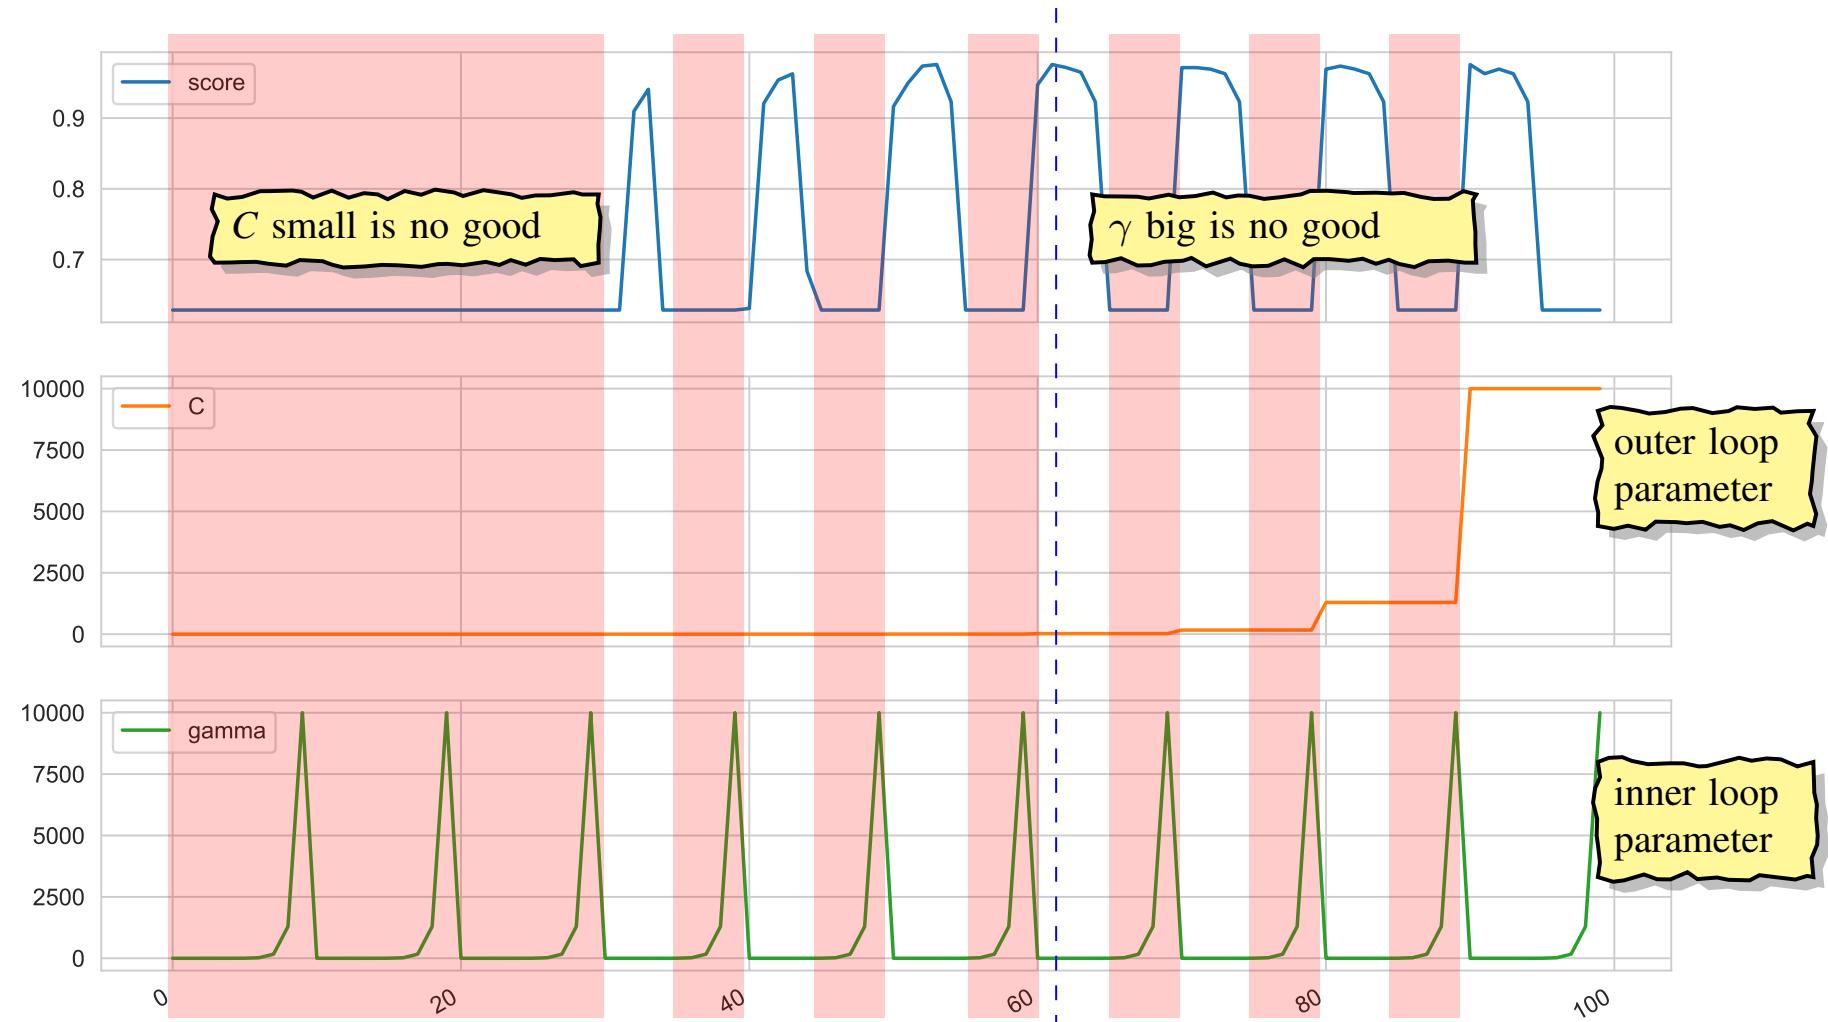
0.9757971014492753  
{'clf\_\_C': 2.782559402207126, 'clf\_\_gamma': 0.046415888336127774, 'clf\_\_kernel': 'rbf'}

- CV score improved from 0.971 with default parameter values to 0.978.
- GridSearchCV with option `return_train_score=True`, returns results of each combination — determine effect of hyperparameters to refine grid search.

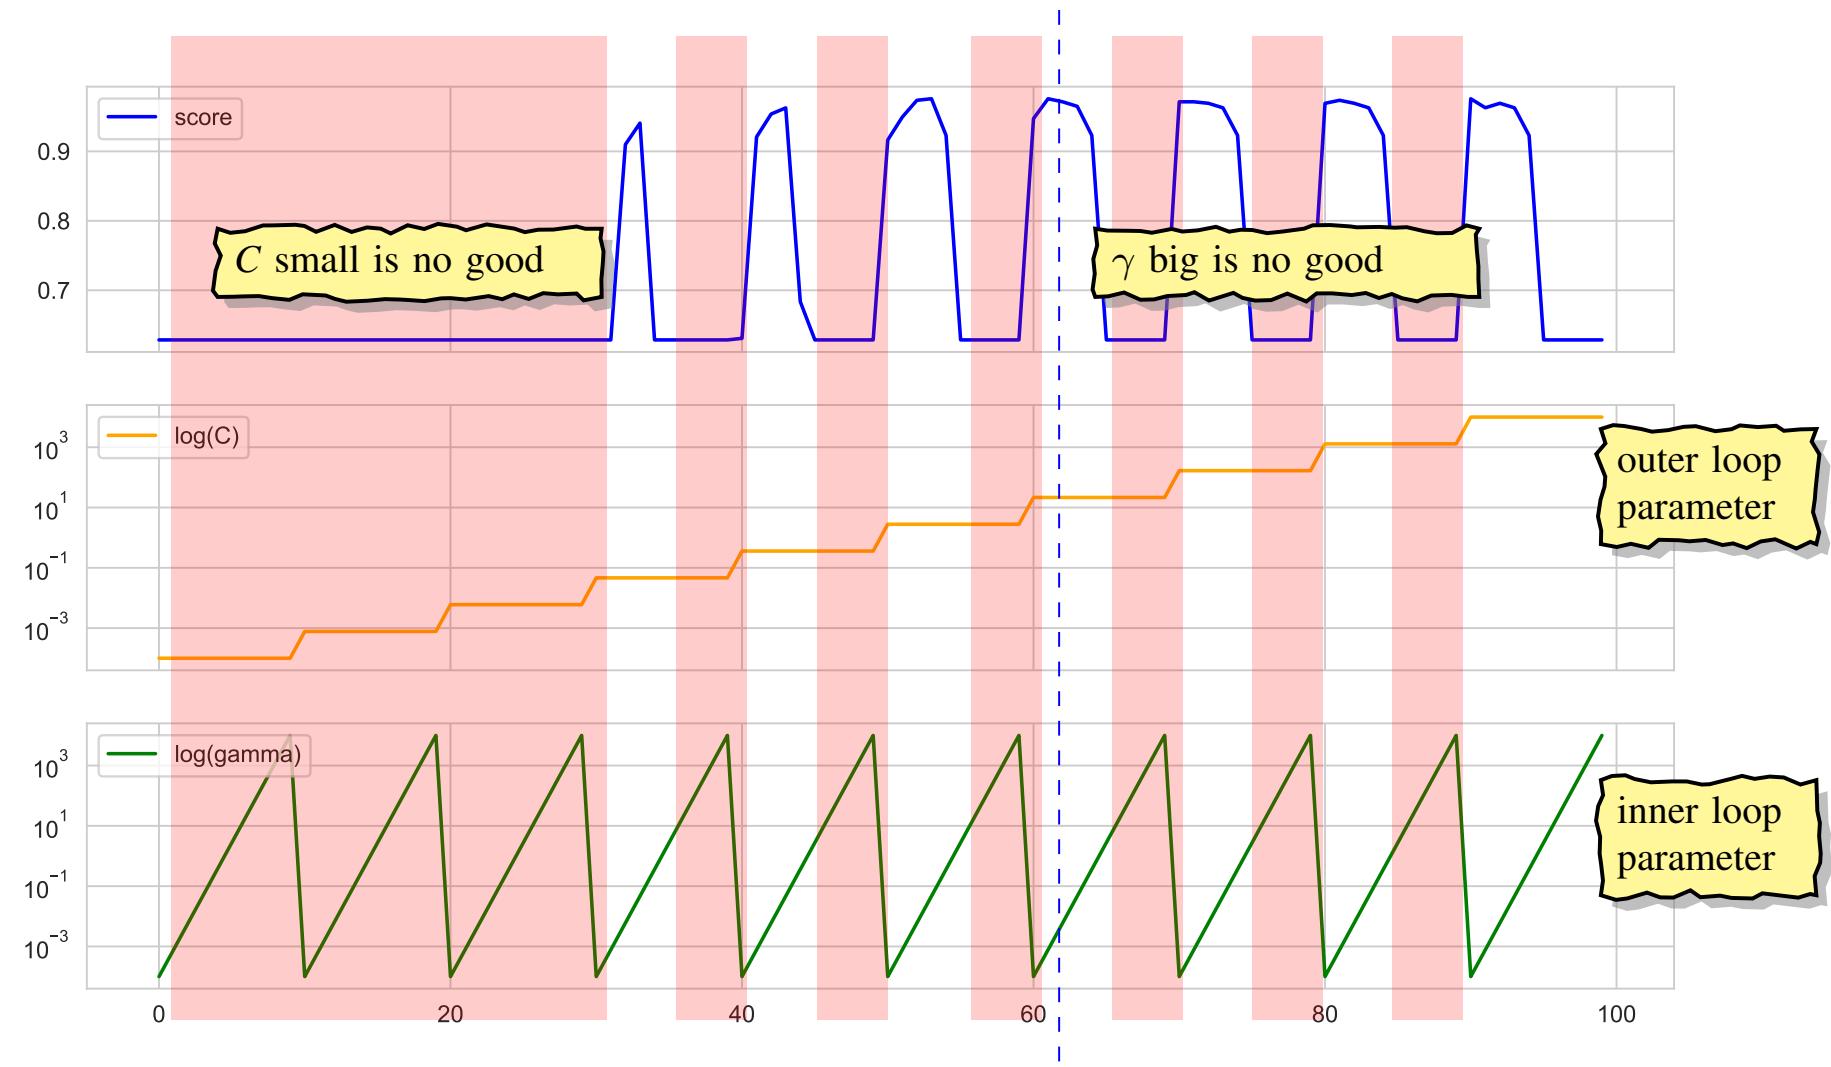
3  
3. df\_gs = pd.DataFrame(np.transpose([  
 gs.cv\_results\_[ "mean\_test\_score" ],  
 gs.cv\_results\_[ "param\_clf\_C" ].data,  
 gs.cv\_results\_[ "param\_clf\_gamma" ].data]),  
 columns=[ 'score', 'C', 'gamma' ])

df\_gs.plot(subplots=True, figsize=(12, 8))  
plt.savefig("gs\_\_svm\_\_C\_gamma.pdf", bbox\_inches="tight")  
plt.show()

# Grid Search Example 1 ( $C = 10, \gamma = 0.01$ )



# Grid Search Example 1 ( $C = 10$ , $\gamma = 0.01$ ) (Using log scales)



## Grid Search Example 2

GridSearch can accept a list of dictionaries so that incompatible hyperparameter combinations can still be searched. For example, different SVM kernels have different parameters:

```
4 pipeline = Pipeline([
    ('scl', StandardScaler()),
    ('clf', SVC(random_state=SEED))
])

from sklearn.model_selection import GridSearchCV
param_range = [0.0001, 0.001, 0.01, 0.1, 1.0, 10.0, 100.0, 1000.0]
param_grid = [
    {
        'clf__C': param_range,
        'clf__kernel': ['linear']
    },
    {
        'clf__gamma': param_range,
        'clf__C': param_range,
        'clf__kernel': ['rbf']
    }
]
```

0.9780676328502415  
 {'clf\_\_C': 545.5594781168514, 'clf\_\_gamma': 0.0001}

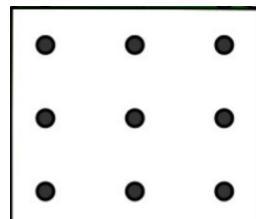
- How many hyperparameter combinations are now generated?

# Random Search

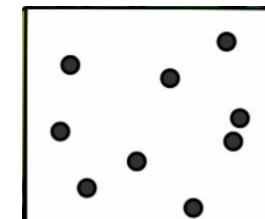
## Random Search

Try random hyperparameter combinations and choose the one with the best cross validation results.

Grid Search



Random Search



### Advantages

- On every step random search varies all parameters  $\Rightarrow$  doesn't spend time on meaningless parameters.
- On average finds near optimal parameters much faster than Grid search.
- It is not limited by grid when optimising continuous parameters.

### Disadvantages

- It may not find the global optimal parameter on a grid.
- All steps are independent. Does not use information about the results gathered to inform search.

# Random Search

```
5 pipeline = Pipeline([
    ('scl', StandardScaler()),
    ('clf', SVC(kernel='rbf', random_state=SEED))
])

from sklearn.model_selection import RandomizedSearchCV
from scipy.stats import randint

param_range = np.logspace(-4, 4, 20)

param_grid = {'clf__C': param_range,
              'clf__gamma': param_range}

rs = RandomizedSearchCV(estimator=pipeline,
                        param_distributions=param_grid,
                        n_iter=100, random_state=SEED,
                        return_train_score=True,
                        scoring='accuracy', cv=10, n_jobs=-1)
```

- This is the same setup as used in GridSearch but doubled the number of parameter values in both  $C$  and  $\gamma$ , so search space is 400 points (not 100). Sampling 100 points.

# Random Search Example

To perform grid search, we call the `fit` method as usual ...

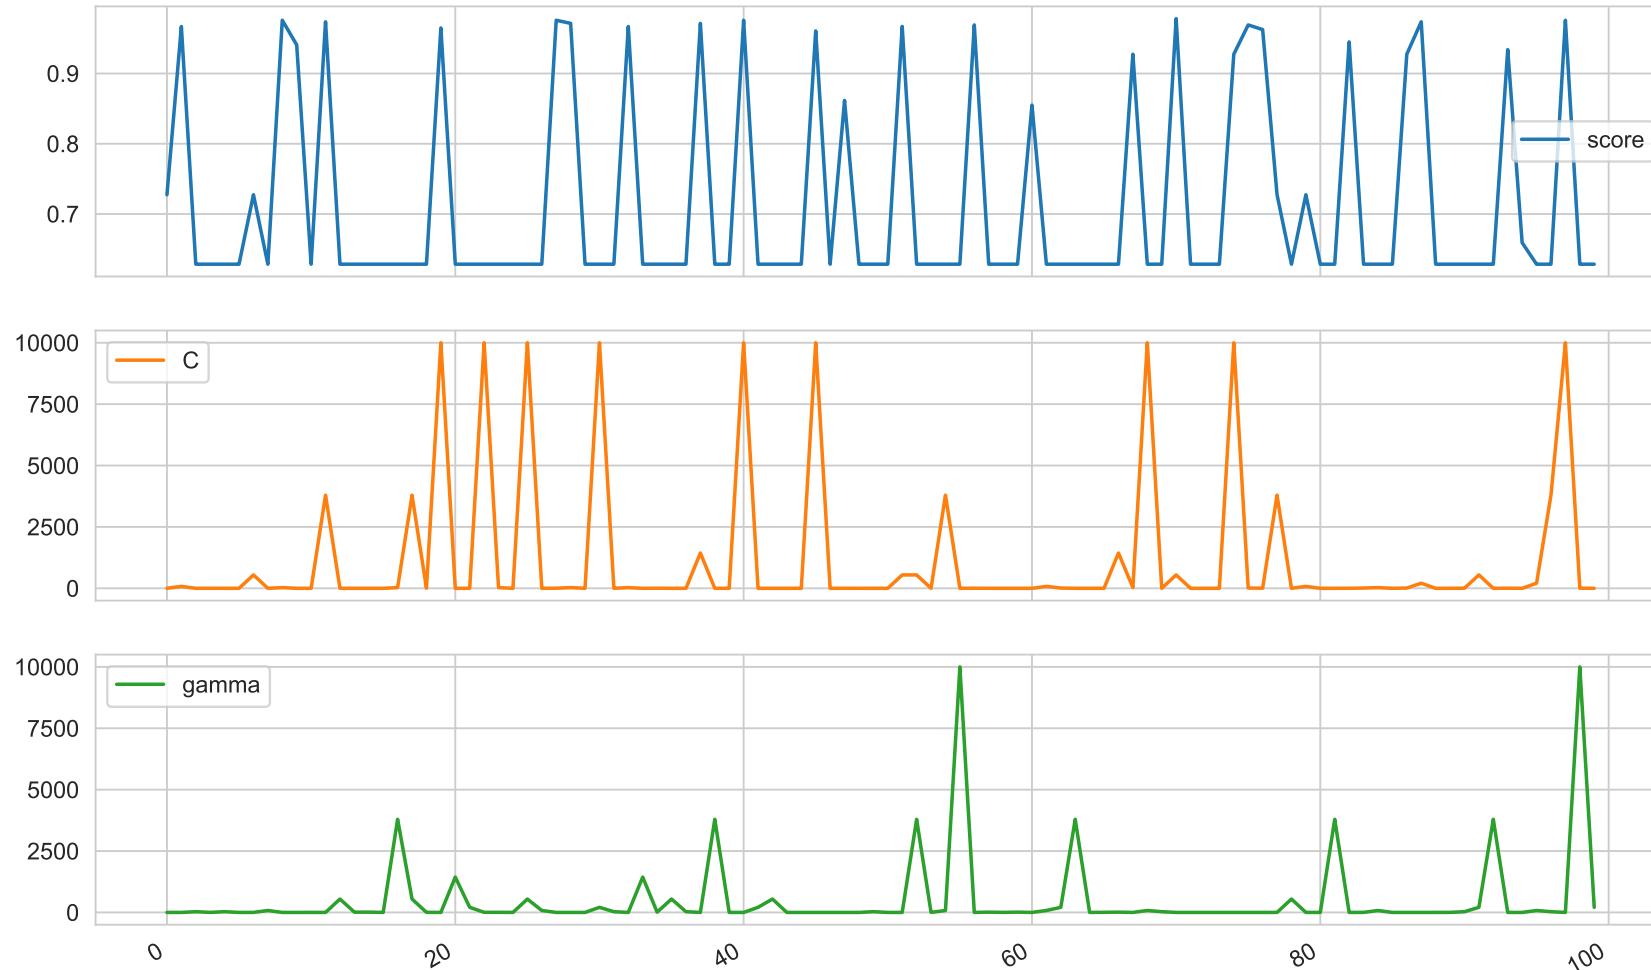
6  
`rs.fit(X_train, y_train)  
print(rs.best_score_)  
print(rs.best_params_)`

```
0.9780676328502415  
{'clf__gamma': 0.0001, 'clf__C': 545.5594781168514}
```

- Sampling half the number of points (50) in a search space 4 times larger, RandomSearch outperforms GridSearch.

7  
`df_rs = pd.DataFrame(np.transpose([  
 rs.cv_results_["mean_test_score"],  
 rs.cv_results_["param_clf_C"].data,  
 rs.cv_results_["param_clf_gamma"].data]),  
 columns=['score', 'C', 'gamma'])  
  
df_rs.plot(subplots=True, figsize=(12, 8))  
plt.savefig("rs__example_1.pdf", bbox_inches="tight")  
plt.show()`

# Random Search Example ( $C = 10000, \gamma = 0.0002636$ )



Every step is completely random.  $\Rightarrow$  does not spend time on useless parameters, but does not use the information gathered on the first steps to improve outcomes of the latter ones.

# The Problem

Both GridSearch and RandomSearch ignore information gained during the search, by using such information can searches improve?

- What information would be useful?
  - Function value
  - Function derivative (gradient) — gives direction in which the function is decreasing (locally).
- Why should we do this?
  - It seems reasonable that lower function values would be clustered — so focus search near best function values found to date.
  - Making multiple decisions based on local information could locate global minimum.
- What can go wrong?
  - Multiple local minimum — can get stuck in a side valley.
    - There will be situations in which we should make decisions against current function value information — see Simulated Annealing.
  - Gradient could be flat (so no direction to go in), or worse almost flat (some algorithms become unstable), or non-existent.

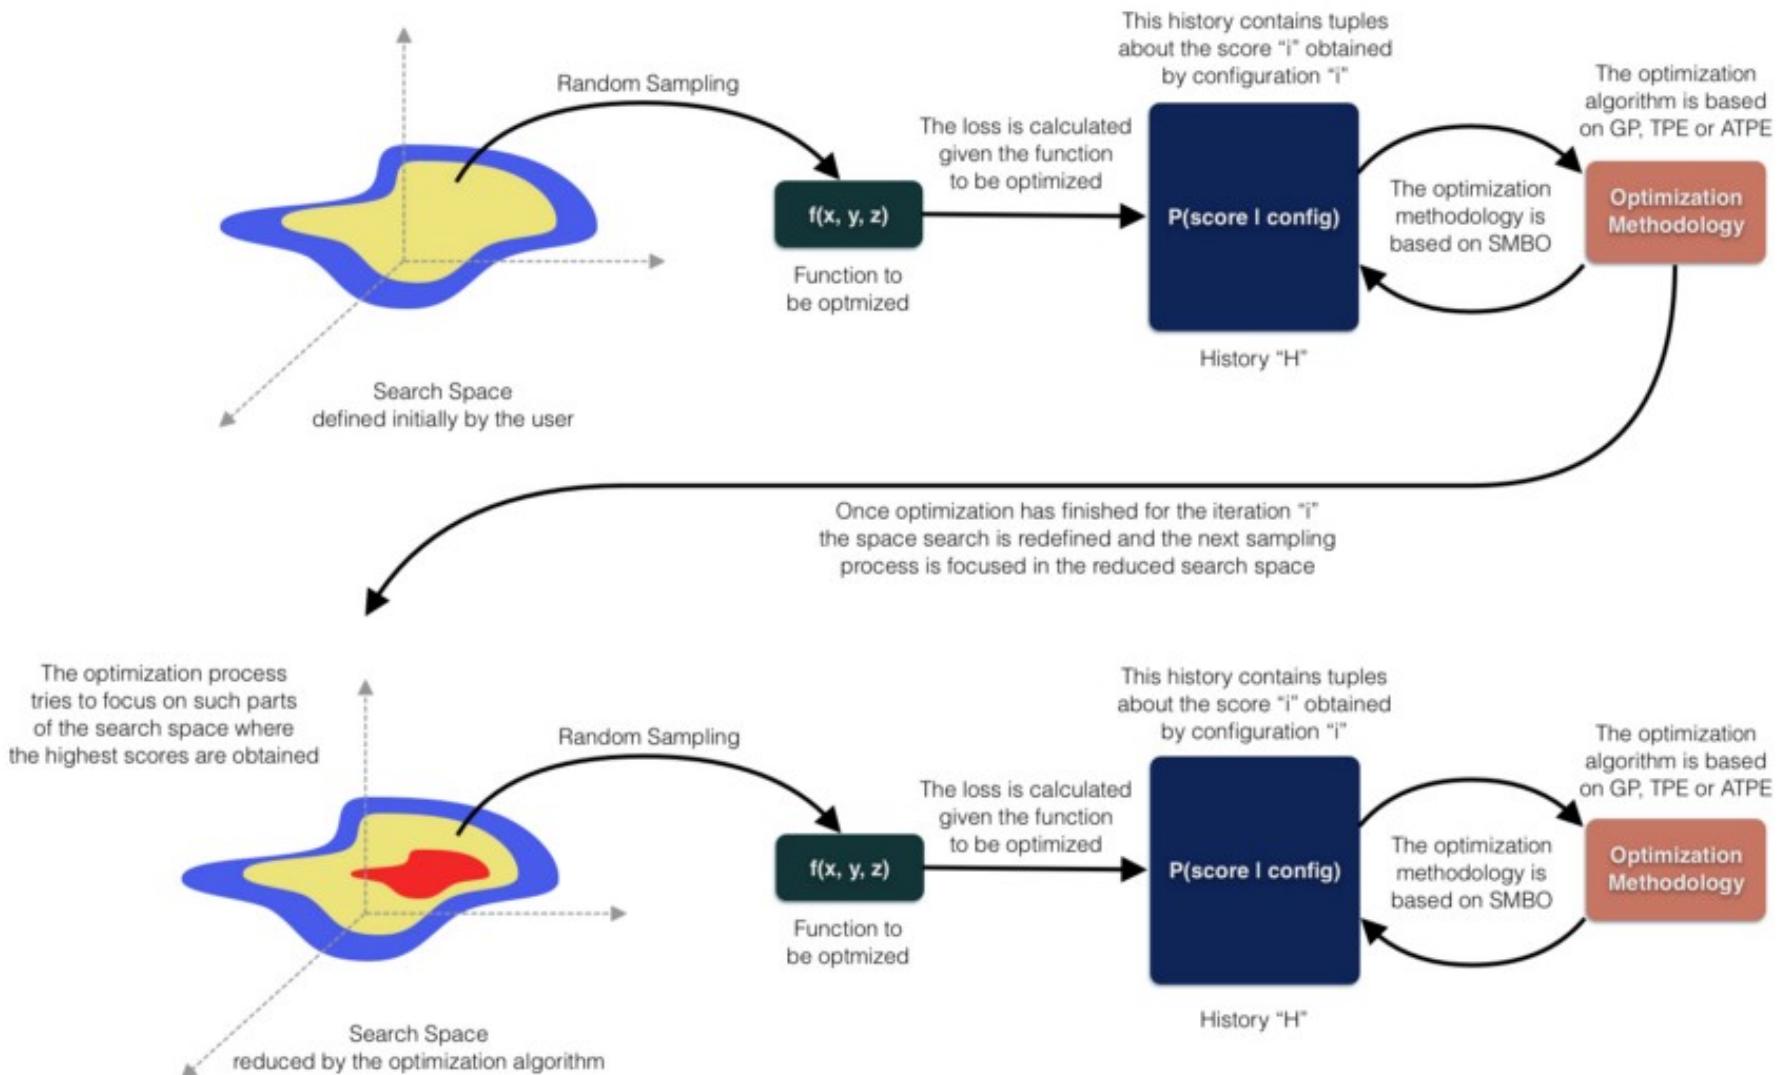
# Bayesian Optimisation

## Sequential model-based optimisation (SMBO)

(also known as **Bayesian optimisation**) is a general technique for function optimisation that uses information from earlier function evaluations to refine the optimisation search — with the aim of minimising number of function calls.

- For sufficiently complex functions, SMBO are some of the most call-efficient (in terms of function evaluations) optimisation methods currently available.
- Compared with standard optimisation strategies such as conjugate gradient descent methods, model-based optimisation algorithms invest more time between function evaluations in order to reduce the number of function evaluations overall.
- Takes advantage of smoothness without analytic gradient.
- Supports real-valued, discrete, and conditional variables.
- Performs relatively well in high dimensions — hundreds of variables, even with budget of just a few hundred function evaluations, i.e., number of function calls is  $\mathcal{O}(\text{dimensions})$  not  $\mathcal{O}(e^{\text{dimensions}})$ .

# Sequential model-based optimisation (in HyperOpt)



# Hyperopt

## Hyperopt: Distributed Asynchronous Hyper-parameter Optimisation

- Python library for optimising over awkward search spaces with real-valued, discrete, and conditional dimensions.

[github.com/hyperopt/hyperopt](https://github.com/hyperopt/hyperopt)

- Install (Anaconda) or use pip  
`conda install -cy conda-forge hyperopt`
- Supports parallel evaluation
  - Uses MongoDB to share results of functions evaluations..

## hyperopt-sklearn

- Hyperopt-based model selection algorithms in scikit-learn

[github.com/hyperopt/hyperopt-sklearn](https://github.com/hyperopt/hyperopt-sklearn)

- Install using pip.

Purpose of hyperopt is not to always to completely optimise the search space but simply to do better than people.

— J Bergstra, SciPy 2013

# Resources

- **Tune Hyperparameters for Classification Machine Learning Algorithms**

[machinelearningmastery.com/](http://machinelearningmastery.com/hyperparameters-for-classification-machine-learning-algorithms/)

[hyperparameters-for-classification-machine-learning-algorithms/](http://hyperparameters-for-classification-machine-learning-algorithms/)

List of important hyper-parameters for sklearn classifiers

- **Hyperopt: A Python Library for Optimizing the Hyperparameters of Machine Learning Algorithms**

[conference.scipy.org/proceedings/scipy2013/pdfs/bergstra\\_hyperopt.pdf](http://conference.scipy.org/proceedings/scipy2013/pdfs/bergstra_hyperopt.pdf)

SCIPY Conference 2013 Paper

- **A Conceptual Explanation of Bayesian Hyperparameter Optimization for Machine Learning**

[towardsdatascience.com/](http://towardsdatascience.com/a-conceptual-explanation-of-bayesian-model-based-hyperparameter-optimization-for-machine-learning-1f3a2a2a2a2a)

[a-conceptual-explanation-of-bayesian-model-based-hyperparameter-optimization-for-machine-learning-1f3a2a2a2a2a](http://a-conceptual-explanation-of-bayesian-model-based-hyperparameter-optimization-for-machine-learning-1f3a2a2a2a2a)

Relatively high level comparison of hyper-parameter tuning approaches.